

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv('dataset.csv')
```

```
df.sample(5)
```

↗

	Unnamed: 0	track_id	artists	album_name	track_name	popularity
69866	69866	3hW6C2zvNurb9gIR3tckrV	K. G. Markose	Marian	Jeevanekidum	
83100	83100	0b6wdul3A5sQNpIOv03OxP	Duke Dumont	Ocean Drive	Ocean Drive	
60645	60645	29gFw3PNupQzPv0XiYrAXm	Intence	Leave If Uh Waah Leave	Leave If Uh Waah Leave	
44710	44710	5jsDxDkJ1PqyYUWhDMr86B	Stone Sour	Hydrograd	Song #3	
4160	4160	7gbSDq9luQx6yVI7HyJGIW	Peter Sandberg	Dismantle	Dismantle	

5 rows × 21 columns

```
df.shape
```

↗

(114000, 21)

```
df.isnull().value_counts()
```

↗

Unnamed: 0	track_id	artists	album_name	track_name	popularity	duration_ms	explicit
False	False	False	False	False	False	False	False
		True	True	True	False	False	False

dtype: int64

```
df.columns
```

```
Index(['Unnamed: 0', 'track_id', 'artists', 'album_name', 'track_name',
      'popularity', 'duration_ms', 'explicit', 'danceability', 'energy',
      'key', 'loudness', 'mode', 'speechiness', 'acousticness',
      'instrumentalness', 'liveness', 'valence', 'tempo', 'time_signature',
      'track_genre'],
      dtype='object')
```

```
df.drop(['Unnamed: 0'],axis = 1,inplace = True)
df.sample(5)
```

		track_id	artists	album_name	track_name	popularity
19251	57buRfUBYm7fFFoIM78qbs	Kacey Musgraves	Best Alternative Pop Tunes	easier said	0	
55359	4YwGnKgtq7V0EI1hYNOirm	Pritam;Shreya Ghoshal	Action Replay	O Bekhabar	52	
104445	1d6B9SauXraTqsUULNIGxD	Ana Torroja;Sentidos Opuestos	90's Pop Tour 4 (En Vivo Desde la Arena Ciudad...	Hijo de la Luna (En Vivo)	31	
47201	225xvV8r1yKMHErSWivnow	Aerosmith	Armageddon - The Album	I Don't Want to Miss a Thing - From "Armageddo...	74	
75651	4r48ds35jyBVxzK4gS1NLW	Himekami	Sennen Kairou	Sennen no Inori	29	

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 114000 entries, 0 to 113999
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   track_id              114000 non-null object
1   artists               113999 non-null object
2   album_name            113999 non-null object
3   track_name            113999 non-null object
4   popularity            114000 non-null int64
5   duration_ms           114000 non-null int64
6   explicit              114000 non-null bool
7   danceability          114000 non-null float64
8   energy                114000 non-null float64
9   key                   114000 non-null int64
..  ..

```

```

10 loudness          114000 non-null float64
11 mode              114000 non-null int64
12 speechiness       114000 non-null float64
13 acousticness       114000 non-null float64
14 instrumentalness   114000 non-null float64
15 liveness           114000 non-null float64
16 valence            114000 non-null float64
17 tempo              114000 non-null float64
18 time_signature     114000 non-null int64
19 track_genre        114000 non-null object
dtypes: bool(1), float64(9), int64(5), object(5)
memory usage: 16.6+ MB

```

```
df.describe()
```

	popularity	duration_ms	danceability	energy	key	lc
count	114000.000000	1.140000e+05	114000.000000	114000.000000	114000.000000	114000.
mean	33.238535	2.280292e+05	0.566800	0.641383	5.309140	-8.
std	22.305078	1.072977e+05	0.173542	0.251529	3.559987	5.
min	0.000000	0.000000e+00	0.000000	0.000000	0.000000	-49.
25%	17.000000	1.740660e+05	0.456000	0.472000	2.000000	-10.
50%	35.000000	2.129060e+05	0.580000	0.685000	5.000000	-7.
75%	50.000000	2.615060e+05	0.695000	0.854000	8.000000	-5.
max	100.000000	5.237295e+06	0.985000	1.000000	11.000000	4.

Minimum value of duration is 0, as seen from df.describe().

This is very peculiar => let's have a look at such tracks, they are probably disposable.

```
df.loc[df['duration_ms'] == 0]
```

	track_id	artists	album_name	track_name	popularity	duration_ms
65900	1kR4glb7nGxHPI3D2ifs59	NaN	NaN	NaN	0	0

The above track is the only one that has any missing values => we will remove it from the dataset.

```

df.drop(df.loc[df['duration_ms'] == 0].index, inplace = True)
df.isnull().sum()

```

	0
track_id	0
artists	0
album_name	0
track_name	0
popularity	0
duration_ms	0
explicit	0
danceability	0
energy	0
key	0
loudness	0
mode	0
speechiness	0
acousticness	0
instrumentalness	0
liveness	0
valence	0
tempo	0
time_signature	0
track_genre	0

dtype: int64

Creating a subset df containing features that pertain to the technical aspects of a song, suspect it might come in handy later.

```
technicalities = df[['track_id', 'acousticness', 'danceability', 'energy', 'key', 'loudness'],
technicalities.sample(5)
```

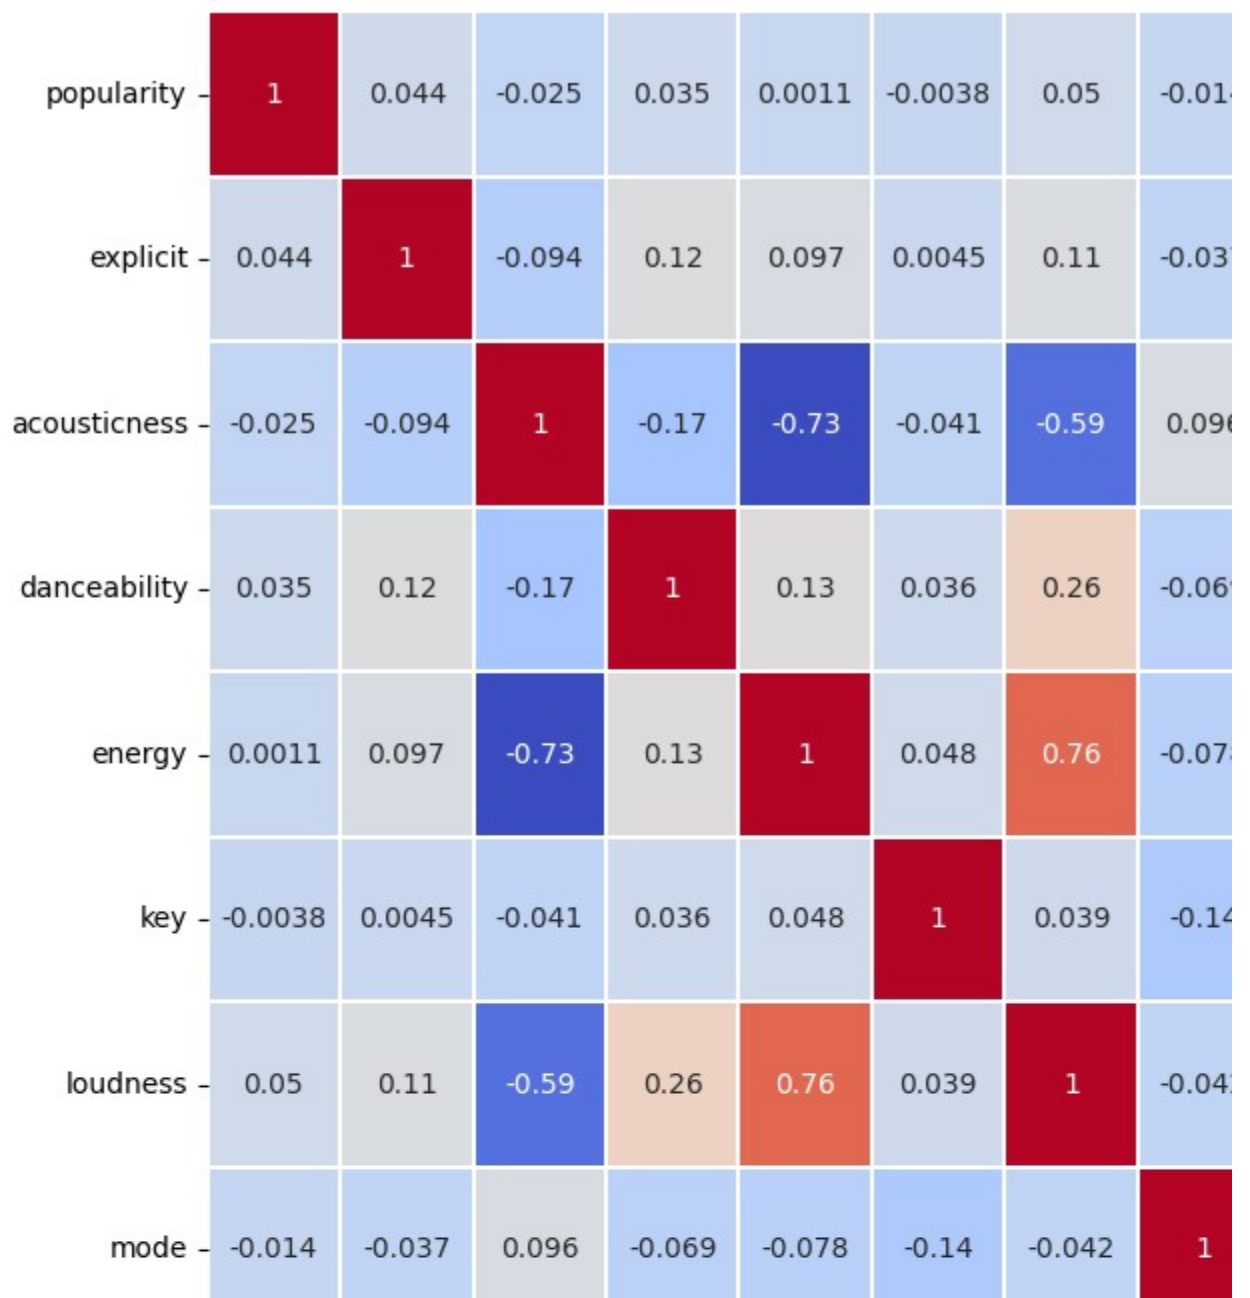
	track_id	acousticness	danceability	energy	key	loudness
62279	0YAMQSmHk6BSUGTYpaoqTJ	0.0286	0.379	0.690	4	-4.790
57540	7Mh1... NYDZ...	0.0550	0.044	0.000	0	0.000


57548	/M3bjqyNXPKL3IHZVJRNdU	0.8550	0.314	0.309	2	-3.809
67733	1ObXAJokXug3fmwrVtlj7v	0.1940	0.737	0.607	0	-3.599
59477	4j6qfCtWcSVUls9zmWHuNI	0.0909	0.658	0.549	1	-17.113
113934	73Elpn5AcedCWPOmxWMxH9	0.0925	0.317	0.562	11	-8.719

✕ Correlation Matrix

```
fig, ax = plt.subplots(figsize=(15,15))
sns.heatmap(df[['popularity','explicit','acousticness','danceability','energy','key', 'loudness', 'mode']])
```

<Axes: >





OBSERVATIONS FROM CORRELATION MATRIX

1. Acousticness vs Energy

Relatively Strong Negative correlation (-0.73) between "Acousticness" and "Energy" => Acoustic music is generally more laid back and less energetic.

2. Acousticness vs Loudness

Moderate Negative (-0.59) => Acoustic music is generally quieter.

3. Danceability and Valence

Moderate Positive Correlation (0.48)

4. Loudness and Energy

--

Moderate to Strong Positive correlation (0.78), louder music is more energetic.

5. Speechiness and Explicit

Weak to Moderate correlation (0.31) => Music containing explicit verses tends to have more verbose lyrics.



✓ Mode column


```
df['mode'].values    #only 1s and 0s
df['mode'].nunique()
```

2

Mode has only two types of values 1s and 0s (likely major and minor modes)

```
df['mode'].value_counts().plot(kind="pie", autopct = "%.2f", figsize=(7,7))
```

<Axes: ylabel='count'>





▼ Duration

```
df['duration_m'] = df['duration_ms']/60000    #converting ms to min and eventually droppi  
df.drop(['duration_ms'],axis = 1,inplace = True)  
  
df['duration_m']
```

	duration_m
0	3.844433
1	2.493500
2	3.513767
3	3.365550
4	3.314217
...	...
113995	6.416650
113996	6.416667
113997	4.524433
113998	4.731550
113999	4.030433

113999 rows × 1 columns

dtype: float64

```
df['duration_m'].skew()  
  
np.float64(11.195825995953085)
```

```
df['duration_m'].median()
```



```
3.5484333333333336
```

```
df['duration_m'].mode()
```

	duration_m
0	2.71495

```
dtype: float64
```

We see that the median is greater than the mode => this data is **positively skewed** (this is also evident from the following graphs)

```
sns.distplot(df['duration_m'])
```

```
<ipython-input-21-add6c3771088>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['duration_m'])  
<Axes: xlabel='duration_m', ylabel='Density'>
```



0 20 40 60 80
duration_m

```
sns.boxplot(df['duration_m'])
```

<Axes: ylabel='duration_m'>



Both the distplot and boxplot tell us that MOST songs have a short duration (median = 3.54 min). However there are also several outliers. We have used the IQR method to classify points as outliers, so any duration longer than $(1.5 \times \text{IQR} + Q3)$ has been branded as outlier. We can see that 5344 entries satisfy this condition => 4.6877% of total entries can be considered as outliers.

```
df['duration_m'].describe()
```

	duration_m
count	113999.000000
mean	3.800519
std	1.788268
min	0.143100
25%	2.901100
50%	3.548433
75%	4.250122

```
15%      4.358433
```

```
max      87.288250
```

```
dtype: float64
```

```
durationIQR = (4.358433 - 2.901100)
```

```
outliersCounts = df.loc[df['duration_m'] >= (1.5*(durationIQR) + 4.358433)].value_counts()
```

```
print(f"Percentage of Outliers = {100 * outliersCounts/df.shape[0]}")
```

```
Percentage of Outliers = 4.687760418951044
```

The duration data is heavily right skewed as determined earlier. This means that we can apply the *LOG TRANSFORM* during the feature engineering phase on the Duration column.

```
transfDuration = np.log(df['duration_m'])
```

```
sns.distplot(transfDuration)
```

```
<ipython-input-25-aec7c8f53997>:2: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(transfDuration)
<Axes: xlabel='duration_m', ylabel='Density'>
```





```
transfDuration.skew()
```

```
np.float64(-0.31831854307791757)
```

```
df['duration_m'] = transfDuration
```


We can see that the LOG Transform has brought the data to closer to being a normal distribution. This is confirmed visually thru the graph as well as mathematically since the skew value of the transformed data is now much closer to 0.


✓ Time Signature

Analysing common Time Signatures, we see from the following graphs that 4/4 is by far the most commonly used time signature.

```
plt.bar(df['time_signature'].unique(),df['time_signature'].value_counts())
```


<BarContainer object of 5 artists>





```
df['time_signature'].value_counts().plot(kind = 'pie', autopct = "%2.f", figsize = (7,7))
```

```
<Axes: ylabel='count'>
```



✓ A Hunch about Time Signatures and Subsequent Testing



From the above results we can see that "4" vastly outnumbers the other time signatures. I feel it would make sense to turn this feature into a binary feature of "4 vs Not-4".

At this point, I have a hunch that songs with **time signature 4** may be **more popular** than songs with other time signatures or vice-versa. We will test this hunch now.

```
binaryTimeSig = df['time_signature'].apply(lambda x: 1 if x == 4 else 0)

df2 = pd.DataFrame({'Popularity':df['popularity'],'binaryTimeSig':binaryTimeSig})

df2.corr()
```

	Popularity	binaryTimeSig	
Popularity	1.000000	0.056315	
binaryTimeSig	0.056315	1.000000	

A correlation of **0.056315** is considered **very weak**. *Therefore my hunch was incorrect.*

We can conclude that in this dataset, there is little to no correlation between whether a song is in 4/4 time signature or not.

```
df2[df2['binaryTimeSig']==1]['Popularity'].describe()
```

	Popularity
count	101842.000000
mean	33.672807
std	22.378415
min	0.000000
25%	17.000000
50%	35.000000
75%	50.000000
max	100.000000

dtype: float64

▼ TEMPO

Analysing Tempo Column

```
df['tempo'].describe()
```

tempo

count	113999.000000
mean	122.147695
std	29.978290
min	0.000000
25%	99.218500
50%	122.017000
75%	140.071000
max	243.372000

dtype: float64

```
sns.distplot(df['tempo'],bins = 20)
```


<ipython-input-33-06e52430cf90>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['tempo'],bins = 20)
<Axes: xlabel='tempo', ylabel='Density'>
```





✓ Popularity

Analysing the Popularity Column

```
df['popularity'].describe()
```

	popularity
count	113999.000000
mean	33.238827
std	22.304959
min	0.000000
25%	17.000000
50%	35.000000
75%	50.000000
max	100.000000

dtype: float64

```
sns.distplot(df['popularity'])
```

<ipython-input-35-93de2ca8ddf5>:1: UserWarning:



`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['popularity'])
<Axes: xlabel='popularity', ylabel='Density'>
```





▼ Just for Reference

df.columns

```
Index(['track_id', 'artists', 'album_name', 'track_name', 'popularity',
      'explicit', 'danceability', 'energy', 'key', 'loudness', 'mode',
      'speechiness', 'acousticness', 'instrumentalness', 'liveness',
      'valence', 'tempo', 'time_signature', 'track_genre', 'duration_m'],
      dtype='object')
```

df['track_genre'].value_counts()

	count
track_genre	
acoustic	1000
afrobeat	1000
alt-rock	1000
alternative	1000
ambient	1000
...	...
techno	1000

turkish	1000
trip-hop	1000
world-music	1000
k-pop	999

114 rows × 1 columns

dtype: int64

✓ Groupby Genre

```
genreWise = pd.DataFrame(df.groupby("track_genre"))
```

✓ Attempting Clustering of Using Kmeans

This dataset has 114 genres in it, most of which are very granular and specific. We will attempt to group together songs with similar features into clusters using kmeans algorithm.


```
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

```
features = df.drop(['track_genre','track_id', 'artists', 'album_name', 'track_name'], axi
```

```
scaler = StandardScaler()
features_scaled = scaler.fit_transform(features)
```

```
inertia = []
K = range(1, 35)
for k in K:
    kmeans = KMeans(n_clusters=k, random_state=45)
    kmeans.fit(features_scaled)
    inertia.append(kmeans.inertia_)
```

```
plt.plot(K, inertia, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method For Optimal k')
plt.show()
```



```
numClust = 7
```

```
kmeans = KMeans(n_clusters=numClust, random_state=42, init='k-means++')
df['cluster'] = kmeans.fit_predict(features_scaled)
```

```
genre_cluster_ct = pd.crosstab(df['cluster'], df['track_genre'])
print(genre_cluster_ct)
```

track_genre \ cluster	acoustic	afrobeat	alt-rock	alternative	ambient	anime
0	10	58	63	31	10	33
1	32	6	1	1	651	157
2	60	314	286	299	11	261
3	608	89	89	97	181	56
4	234	282	463	382	40	311
5	50	19	56	164	4	55
6	6	232	42	26	103	127

track_genre \ cluster	black-metal	bluegrass	blues	brazil	...	spanish	study
0	57	89	37	201	...	94	22
1	18	34	8	2	...	1	254
2	162	75	200	178	...	264	40
3	11	422	341	205	...	145	174
4	154	304	385	321	...	445	22
5	110	5	11	86	...	37	2

5	118	5	14	80	...	51	0
6	480	71	15	7	...	14	488

track_genre	swedish	synth-pop	tango	techno	trance	trip-hop	turkish	\
cluster								
0	34	38	73	22	59	47	30	
1	6	1	54	4	0	18	8	
2	194	308	61	160	228	214	394	
3	258	100	791	22	2	115	184	
4	401	456	9	150	192	169	145	
5	95	44	0	25	24	45	218	
6	12	53	12	617	495	392	21	

track_genre	world-music
cluster	
0	152
1	104
2	56
3	269
4	418
5	0
6	1

[7 rows x 114 columns]

...

```
from sklearn.cluster import DBSCAN
```

```
# Scale your data
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(features)
```

```
# Run DBSCAN
```

```
dbscan = DBSCAN(eps=3, min_samples=10)
```

```
df['cluster'] = dbscan.fit_predict(X_scaled)
```

...

```
\nfrom sklearn.cluster import DBSCAN\n\n# Scale your data\nscaler = StandardScaler\n()\nX_scaled = scaler.fit_transform(features)\n\n# Run DBSCAN\nndbscan = DBSCAN(eps=3\n min samples=10)\ndf['cluster'] = dbscan.fit_predict(X_scaled)\n'
```

```
from sklearn.decomposition import PCA
```

```
pca = PCA(n_components=3)
```

```
components = pca.fit_transform(features_scaled)
```

```
plt.scatter(components[:, 0], components[:, 1], c=df['cluster'], cmap='tab10')
```

```
plt.xlabel('PC1')
```



```
plt.ylabel('PC2')
```


```
plt.title('KMeans Clusters Visualized (PCA)')
```

```
plt.show()
```


```
plt.scatter(components[:, 0], components[:, 2], c=df['cluster'], cmap='tab10')  
plt.xlabel('PC1')  
plt.ylabel('PC3')  
plt.title('KMeans Clusters Visualized (PCA)')  
plt.show()
```

```
plt.scatter(components[:, 1], components[:, 2], c=df['cluster'], cmap='tab10')  
plt.xlabel('PC2')  
plt.ylabel('PC3')  
plt.title('KMeans Clusters Visualized (PCA)')  
plt.show()
```





KMeans Clusters Visualized (PCA)



```
# Create 3D scatter plot
fig = plt.figure(figsize=(10, 7))
ax = fig.add_subplot(111, projection='3d')

scatter = ax.scatter(
    components[:, 0], components[:, 1], components[:, 2],
    c=df['cluster'], cmap='tab10', s=15
)

# Labeling axes and title
ax.set_title("KMeans Clusters Visualized (PCA)", fontsize=12)
ax.set_xlabel("PC1")
ax.set_ylabel("PC2")
ax.set_zlabel("PC3")
```

```
ax.view_init(45,225)

# Add colorbar
fig.colorbar(scatter, ax=ax, label='Cluster')

plt.tight_layout()
plt.show()
```



```
fig = plt.figure(figsize=(10, 7))
```



```
ax = fig.add_subplot(111, projection='3d')

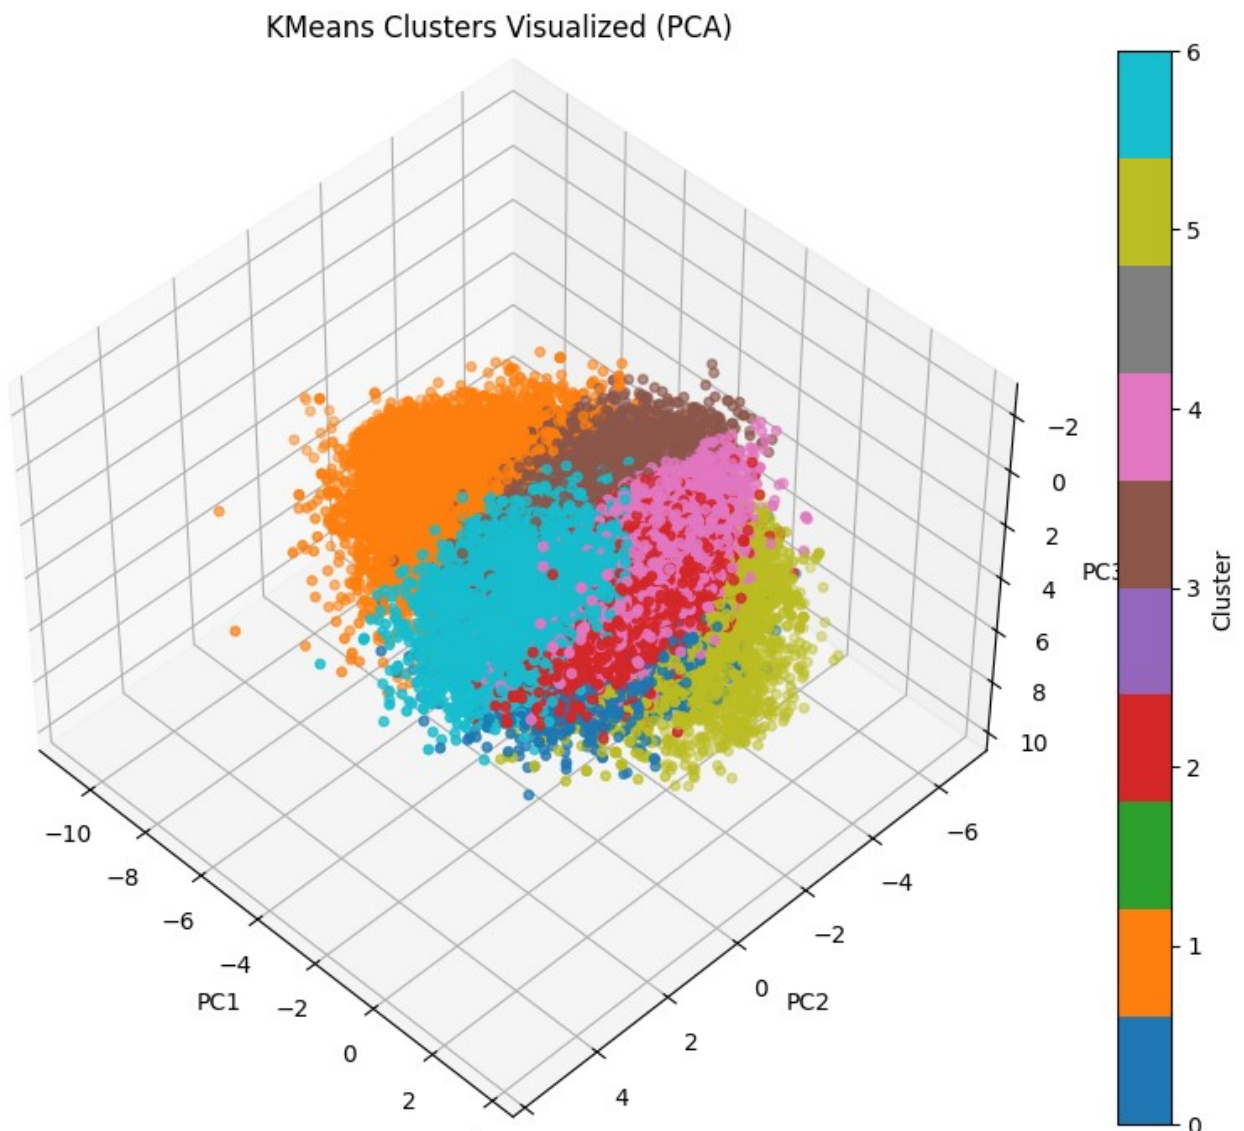
scatter = ax.scatter(
    components[:, 0], components[:, 1], components[:, 2],
    c=df['cluster'], cmap='tab10', s=15
)

ax.set_title("KMeans Clusters Visualized (PCA)", fontsize=12)
ax.set_xlabel("PC1")
ax.set_ylabel("PC2")
ax.set_zlabel("PC3")

ax.view_init(225, 225)

# Add colorbar
fig.colorbar(scatter, ax=ax, label='Cluster')

plt.tight_layout()
plt.show()
```



4

6

```
print(pca.explained_variance_)

[3.00256424 1.5671262 1.38153061]

print(pca.explained_variance_ratio_)

[0.20016919 0.10447416 0.09210123]
```

One can see a fair amount of patterns forming after performing clustering. It is not perfect, we do have some overlap in the clusters, but we must bear in mind that we have only considered 3 Principal Components.

As we can see, the first 3 principal components account for ~20%, ~10%, and ~9.2% respectively, which is not a lot, so it is remarkable that we are still able to see patterns form within the clusters

✓ CLUSTERING OBSERVATIONS

Add blockquote

```
for i in range(numClust):
    print(f"Cluster {i+1}")
    print(df[df['cluster'] == i]['track_genre'].value_counts().head(35))
    print()
```

```
Cluster 1
track_genre
pagode          487
sertanejo       415
samba           409
comedy          276
mpb             271
gospel          222
brazil          201
forro           176
world-music     152
nan             110
```

rock	140
ska	143
heavy-metal	141
sleep	119
happy	109
grindcore	104
groove	99
party	95
spanish	94
hardstyle	93
bluegrass	89
punk-rock	84
hard-rock	83
goth	81
grunge	77
tango	73
psych-rock	71
show-tunes	70
opera	70
rock-n-roll	67
metalcore	64
alt-rock	63
power-pop	62
drum-and-bass	62
punk	61
trance	59

Name: count, dtype: int64

Cluster 2

track_genre	
sleep	792
new-age	772
classical	761
ambient	651
piano	560
guitar	448
disney	440
iranian	437
opera	351
german	299
idm	272
study	254
romance	181
anime	157
british	148
jazz	117
world-music	104

****POSITIVES****

* This clustering certainly isn't perfect, but there are several eye-catching groupings

groupings.

* Cluster 1 has a strong Brazilian theme – Pagode, Sertanejo, Samba, MPB, Forró, and Brazil all show up. RnB and Spanish are also present, which fits, considering musical influences.

* However heavy Metal and Sleep are also present in Cluster 1, which is rather odd and amusing.

* Cluster 2 is full of Ambient, Sleep, New-Age, Piano, Opera, and Study. It also includes Anime, Disney, and Jazz, which makes this one of the most coherent clusters.

* Cluster 3 contains House, EDM, Dance, Hip-Hop, Reggaeton, K-pop, Salsa, and Dubstep. A very energetic mix – lots of rhythm-heavy genres.

* Cluster 4 includes Jazz, Folk, Acoustic, Country, Singer-Songwriter, and Bluegrass. Pretty strong grouping of organic and melodic genres that have shared history, which is remarkable (although "Jazz" is a massive genre and has several sub genres which may or may not be similar to the others in this group).

* Cluster 5 features Punk, Rock, Grunge, J-Rock, J-Pop, Kids, Party, Power-Pop, and Ska. Feels chaotic, but all very youth-focused and upbeat.

* Cluster 6 is kind of aggressive and emotional – Comedy, Emo, Sad, Hardcore, Funk, Grindcore, and Metalcore all show up. A rather peculiar group, perhaps the most inconsistent cluster.

* Cluster 7 is packed with electronic genres – Minimal Techno, Detroit Techno, Techno, Trance, IDM, Trip-Hop, Drum-and-Bass. Also includes Black Metal, Death Metal, and Industrial, which is somewhat odd, but does make some sense sonically (unlike sleep and heavy metal being grouped together in cluster 1).

****NEGATIVES****

* 0.15 silhouette score, which indicates very poor cluster separation.

****FINAL CONCLUSIONS****

Despite a relatively low silhouette score of 0.15, the clustering results reveal meaningful structure in the audio feature space. Several clusters group together genres with shared musical characteristics – such as:

Techno subgenres and house styles

Ambient, study, and sleep tracks

Folk, jazz, and acoustic genres

Overlaps between Latin, Spanish, and RnB

This indicates that even without explicit genre labels, a simple algorithm like KMeans was able to capture some degree of stylistic similarities in the data. However, the presence of genres like "Kids" or "Sleep" in multiple, unrelated clusters highlights the limitations of using purely audio-based features for genre classification.

Ultimately, the 0.15 silhouette score reflects the complexity and fuzziness of real-world music genre boundaries. While the clusters aren't cleanly separable, the model does uncover some valuable patterns – suggesting that unsupervised learning, when combined with domain insight, can still yield practically useful groupings.

POSITIVES

- This clustering certainly isn't perfect, but there are several eye-catching groupings.
- Cluster 1 has a strong Brazilian theme – Pagode, Sertanejo, Samba, MPB, Forró, and Brazil all show up. RNB and Spanish are also present, which fits, considering musical influences.
- However heavy Metal and Sleep are also present in Cluster 1, which is rather odd and amusing.
- Cluster 2 is full of Ambient, Sleep, New-Age, Piano, Opera, and Study. It also includes Anime, Disney, and Jazz, which makes this one of the most coherent clusters.
- Cluster 3 contains House, EDM, Dance, Hip-Hop, Reggaeton, K-pop, Salsa, and Dubstep. A very energetic mix – lots of rhythm-heavy genres.
- Cluster 4 includes Jazz, Folk, Acoustic, Country, Singer-Songwriter, and Bluegrass. Pretty strong grouping of organic and melodic genres that have shared history, which is remarkable (although "Jazz" is a massive genre and has several sub genres which may or may not be similar to the others in this group).
- Cluster 5 features Punk, Rock, Grunge, J-Rock, J-Pop, Kids, Party, Power-Pop, and Ska. Feels chaotic, but all very youth-focused and upbeat.
- Cluster 6 is kind of aggressive and emotional – Comedy, Emo, Sad, Hardcore, Funk, Grindcore, and Metalcore all show up. A rather peculiar group, perhaps the most inconsistent cluster.
- Cluster 7 is packed with electronic genres – Minimal Techno, Detroit Techno, Techno, Trance, IDM, Trip-Hop, Drum-and-Bass. Also includes Black Metal, Death Metal, and Industrial, which is somewhat odd, but does make some sense, especially for the sleep

industrial, which is somewhat odd, but does make some sense sonically (unlike sleep and heavy metal being grouped together in cluster 1).

NEGATIVES

- 0.15 silhouette score, which indicates very poor cluster separation.

FINAL CONCLUSIONS

Despite a relatively low silhouette score of 0.15, the clustering results reveal meaningful structure in the audio feature space. Several clusters group together genres with shared musical characteristics — such as:

Techno subgenres and house styles

Ambient, study, and sleep tracks

Folk, jazz, and acoustic genres

Overlaps between Latin, Spanish, and RnB

This indicates that even without explicit genre labels, a simple algorithm like KMeans was able to capture some degree of stylistic similarities in the data. However, the presence of genres like "Kids" or "Sleep" in multiple, unrelated clusters highlights the limitations of using purely audio-based features for genre classification.

Ultimately, the 0.15 silhouette score reflects the complexity and fuzziness of real-world music genre boundaries. While the clusters aren't cleanly separable, the model does uncover some valuable patterns — suggesting that unsupervised learning, when combined with domain insight, can still yield practically useful groupings.

Start coding or [generate](#) with AI.

